

Statement of Work — AI-Enabled Future Operating Model Program

Client: Meridian Global Services (MGS)

Program: "Project Horizon" — Design & Implementation of an AI-Enabled Future Operating Model

Document type: Statement of Work (SOW)

Version: 1.0

Date: 2026-06-23

Prepared for: MGS Executive Committee & PMO

Commercial model: Fixed-price phases with time-and-materials change budget

1. Background and Business Case

Meridian Global Services is a multinational business-services firm (11,000 employees across 14 countries) whose operating model has not kept pace with advances in generative and agentic AI. Manual, fragmented processes in finance operations, customer service, HR shared services, and knowledge work are driving high cost-to-serve, inconsistent quality, and slow cycle times.

The Executive Committee has approved an enterprise program to design and stand up an **AI-enabled Future Operating Model (FOM)** that embeds AI into core processes, decisions, and ways of working — while maintaining governance, regulatory compliance, and workforce trust. The program is expected to reduce operating cost in targeted functions by 18–25%, improve first-contact resolution by 30%, and materially shorten key process cycle times within 18 months of go-live.

This SOW defines the scope, deliverables, governance, and commercial basis for the design and first wave of implementation of the Future Operating Model.

2. Vision and Strategic Objectives

Vision: A measurably more productive, compliant, and adaptive enterprise in

which AI augments every core process and decision, governed by clear human accountability.

Strategic objectives:

1. Define a target operating model across people, process, technology, data, and governance that is AI-native rather than AI-bolted-on.

1. Deliver three production "lighthouse" use cases proving value in customer service, finance operations, and internal knowledge work.

1. Establish an enterprise AI governance and Responsible-AI framework that satisfies regulatory and risk requirements.

1. Build the platform, data foundation, and operating capabilities (an AI Center of Enablement) to scale beyond the initial use cases.

1. Reskill and transition the affected workforce with minimal disruption and strong change adoption.

3. Scope

3.1 In Scope

- Current-state assessment of the four target functions (finance operations, customer service, HR shared services, knowledge work) and their processes, data, systems, roles, and controls.

- Target Operating Model design across five dimensions: People & Org, Process, Technology & Architecture, Data, and Governance & Controls.

- Design and production deployment of **three lighthouse use cases**:
(a) an AI customer-service assistant integrated with the CRM and knowledge base; (b) an intelligent finance-operations workflow (invoice and exception handling); (c) an enterprise knowledge-assistant over approved internal content.

- Enterprise AI platform foundations: model access/governance, retrieval infrastructure, prompt/agent management, evaluation harness, and monitoring.

- Data foundation work required by the lighthouse use cases (sourcing, quality,

access controls, lineage for in-scope domains).

- AI governance, Responsible-AI, security, and compliance framework and controls, including human-in-the-loop and audit requirements.
- Workforce impact assessment, role redesign, reskilling pathways, and change management for affected teams.
- Operating-model run capability: the AI Center of Enablement (people, process, tooling) to operate and scale post-go-live.

3.2 Out of Scope

- Replacement of the core ERP or CRM platforms (integration only).
- Enterprise-wide data-warehouse re-platforming beyond the in-scope domains.
- Functions and geographies outside the four named functions for this wave.
- Net-new physical infrastructure / data-center build (cloud-first assumed).
- Marketing and product R&D use cases (candidate for a later wave).

4. Target Operating Model — Design Dimensions

- **People & Organization:** future role definitions, AI-augmented job designs, spans and layers, the AI Center of Enablement operating structure, and a RACI for AI decisions and accountability.
- **Process:** redesigned end-to-end processes for the target functions with AI and human-in-the-loop decision points, exception handling, and controls.
- **Technology & Architecture:** reference architecture for the AI platform, integration patterns with ERP/CRM/knowledge systems, and build-vs-buy.
- **Data:** data products, quality standards, access/entitlement model, lineage, and retention required for the use cases and governance.
- **Governance & Controls:** Responsible-AI policy, model risk management, evaluation and monitoring standards, regulatory mapping, and audit.

5. Approach and Workstreams (Phased)

Phase 0 — Mobilize (Weeks 1–3): program setup, governance, baselines, detailed plan, and success metrics.

Phase 1 — Assess & Design (Weeks 4–13): current-state assessment, target

operating model design across the five dimensions, use-case design, and business cases.

Phase 2 — Build Lighthouses & Platform (Weeks 14–30): platform foundations, data work, and iterative build of the three lighthouse use cases with evaluation and security review.

Phase 3 — Pilot & Harden (Weeks 31–40): controlled pilots, Responsible-AI and security sign-off, performance tuning, and operational readiness.

Phase 4 — Go-Live & Transition (Weeks 41–46): production go-live, hypercare, workforce transition, and handover to the AI Center of Enablement.

6. Deliverables

1. Program management plan, integrated schedule, and governance charter.
2. Current-state assessment report for the four functions.
3. Future Operating Model blueprint (all five dimensions) with target RACI.
4. Use-case design specifications and business cases (× 3).
5. AI platform reference architecture and deployed foundation services.
6. Three production lighthouse use cases with evaluation and monitoring.
7. Data foundation deliverables for in-scope domains (catalog, quality, access).
8. Responsible-AI & governance framework, policies, and control set.
9. Security and compliance assessment and sign-off package.
10. Workforce transition plan, reskilling curriculum, and change-adoption plan.
11. AI Center of Enablement operating model and run-books.
12. Final program report, benefits realization baseline, and lessons learned.

7. Milestones and Timeline

The program runs approximately **46 weeks (~11 months)**. Key milestones:

- M1 (Wk 3): Program mobilized; baselines and plan approved.
- M2 (Wk 13): Operating model blueprint and use-case designs approved.
- M3 (Wk 22): AI platform foundation live; first use case in build.
- M4 (Wk 30): All three lighthouse use cases build-complete.
- M5 (Wk 40): Pilots complete; Responsible-AI and security sign-off obtained.
- M6 (Wk 46): Production go-live and transition to the Center of Enablement.

8. Budget and Commercial

Total approved budget: **USD 6.4M**, structured as fixed-price phase payments against milestone acceptance, plus a ****USD 0.6M** time-and-materials change

budget** governed by integrated change control. Cloud and third-party model/licensing costs are estimated at USD 0.9M within the total and are tracked separately against a cost baseline. A management reserve of 8% is held by the sponsor for identified-but-unplanned risk responses.

9. Governance, Roles and Responsibilities

- **Executive Sponsor (MGS COO):** accountable for benefits and funding.
- **Steering Committee:** monthly; approves stage gates and major change.
- **Program Director:** day-to-day accountability for delivery and the plan.
- **Workstream Leads:** Operating Model, Platform/Engineering, Data,

Responsible-AI & Risk, Change & Workforce.

- **AI Governance Board:** approves models, use cases, and Responsible-AI

exceptions; owns the model-risk and audit process.

- **Change Control Board:** assesses all changes for scope, schedule, cost,

quality, and risk impact before approval.

10. Stakeholders

COO (sponsor), CFO, CIO/CTO, Chief Data Officer, Chief Risk & Compliance

Officer, CHRO, function heads (Customer Service, Finance Operations, HR Shared

Services, Knowledge Management), Works Council / employee representatives, Data

Privacy Officer, Internal Audit, and key technology vendors.

11. Assumptions

- Cloud platform and enterprise model access are procured and available.
- Business SMEs from each function are available ~25% FTE for the program.
- Approved internal content for the knowledge assistant is identified and

licensed for such use.

- Existing ERP/CRM expose supported integration APIs.

12. Constraints and Dependencies

- Regulated customer and employee data must remain in-region and comply with

applicable data-protection law; no training on personal data without consent.

- Go-live must precede the next fiscal year-end; no major releases during

quarter-end freezes.

- Workforce changes require Works Council consultation in two jurisdictions.
- Dependency on the parallel CRM upgrade program for one integration point.

13. Initial Risks

- Regulatory / Responsible-AI non-compliance leading to launch delay (high).
- Data quality and access gaps undermining use-case performance (high).
- Workforce resistance and adoption shortfall (medium-high).
- Model performance / hallucination risk in customer-facing use cases (high).
- Vendor and platform lock-in and cost overrun (medium).
- Scope creep across functions beyond the agreed wave (medium).

14. Quality and Acceptance Criteria

- Each deliverable is reviewed against documented acceptance criteria and

signed off by the named approver before milestone payment.

- Lighthouse use cases must meet pre-agreed accuracy, latency, safety, and

business-KPI thresholds in pilot before production sign-off.

- All AI use cases must pass Responsible-AI, security, and privacy review.

15. Compliance, Security and Responsible AI

The program adopts a Responsible-AI framework covering fairness, transparency, human oversight, accountability, security, and privacy. All in-scope models and agents are registered, risk-tiered, evaluated, monitored in production, and auditable. Human-in-the-loop is mandatory for high-impact decisions.

16. Success Metrics (KPIs)

- 18–25% cost-to-serve reduction in targeted functions within 18 months.
- + 30% first-contact resolution in customer service.
- 40% cycle-time reduction in the finance-operations lighthouse.
- ≥ 80% adoption of the knowledge assistant among target users.
- Zero critical Responsible-AI or security incidents at go-live.

End of Statement of Work. This document is the basis for AI Plan Generation and PMO Validation in the PMO Intelligence Engine.